

# ‘votess: A multi-target, GPU-capable, parallel Voronoi tessellator’

Samridh Dev Singh<sup>1</sup>, Chris Byrohl<sup>2</sup>, and Dylan Nelson<sup>2</sup>

<sup>1</sup> Grinnell College, 1115 8th Avenue, 50112 Grinnell, United States of America <sup>2</sup> Heidelberg University, Institute for Theoretical Astronomy, Albert-Ueberle-Str. 2, 69120 Heidelberg, Germany

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Open Journals](#) ↗

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

## License

Authors of papers retain copyright  
and release the work under a  
Creative Commons Attribution 4.0  
International License ([CC BY 4.0](#)).

## Statement of need

A Voronoi tessellation is a spatial decomposition that partitions space into a set of convex hulls based on proximity to a seed points with interesting to the applications in biology, data science, geography, and physics. One compute intensive application is its use in astrophysics, such as the analysis of matter distribution ([Weygaert, 1994](#)), optimal transport theory for early-universe reconstruction ([Levy et al., 2021](#)), and in observational data analysis and numerical simulations of cosmic structure formation ([Springel, 2010](#)).

The increasing size of datasets produced today have underscored the need for more efficient algorithms to both generate and analyse these datasets, and the rise of heterogenous computing facilities would enable such new algorithms to be run. There do exist several sequential and parallel implementations of the Voronoi diagram problem ([Marot et al., 2019](#)) ([Wu et al., 2023](#)) ([The CGAL Project, 2018](#)) ([Inria, 2018](#)), however, they are mostly restricted to CPU or specific GPU architectures, thus limiting their potential as a portable multi-architecture algorithm.


## Summary

votess is a library, implementing Ray’s meshless algorithm ([Ray et al., 2018](#)), for computing parallel three dimensional Voronoi tessellations using the C++/SYCL framework for CPU, GPUs and other future architectures.

One advantage of this algorithm is the ability for each cell to be computed independently ([Ray et al., 2018](#)), making it suitable for parallel execution. It also produces the geometry of the Voronoi cells via their neighbor connectivity information, rather than a full combinatorial mesh data structure, thus making it more ammenable to data parallel architectures than alternatives such as sequential insertion or the Bowyer-Watson algorithm ([Bowyer, 1981](#)) ([Watson, 1981](#)).

The core method of votess consists of two main steps. First, the input set of points is sorted into a grid, and a k-nearest neighbors search is performed. Once the k nearest neighbors are identified for each point, the Voronoi cell is computed by iteratively clipping a bounding box using the perpendicular bisectors between the point and the identified neighbors. A *security radius* condition ([Lévy & Bonneel, 2013](#)) ensures that the resulting Voronoi cell is valid, and if the cell cannot be validated, an CPU fallback mechanism is used.

## 35 Performance



36

37 In Figure 1, we show its performance compared to two other single-threaded Voronoi tessellation  
38 libraries: CGAL and Voro++. Both are well-tested and widely used. CGAL implements Voronoi  
39 Tessellations through a Delaunay mesh using the CPU in parallel ([The CGAL Project, 2018](#)),  
40 while Voro++ computes three-dimensional Voronoi tessellations in a single core via a cell-based  
41 computation approach that is well-suited for physical applications ([Rycroft, 2009](#)). The  
42 benchmark uses a float32 uniformly distributed dataset, as it is currently the simplest to test.  
43 From the results of ([Ray et al., 2018](#)), this dataset would provide the lowest performance, but  
44 higher than clustered datasets: datasets currently beyond the scope of the current version.

45 Other Multithreaded Voronoi tessellation codes exist, including ParVoro++ ([Wu et al., 2023](#)),  
46 and GEOGRAM ([Inria, 2018](#)). However, they do not natively support GPU architectures, and we  
47 are unable to benchmark them.

48 From the graph above, votess outperforms the single-threaded alternative, however, when  
49 compared to the established CGAL, both the CPU and GPU version falls short by around a  
50 factor of 6. Currently the problem is being addressed, and future optimizations should close  
51 the gap.

## 52 Features

53 votess is designed to be versatile. It supports various outputs, including the natural neighbor  
54 information for each Voronoi cell. This is a 2D jagged array of neighbor indices of the sorted  
55 input dataset.

56 Users can invoke votess in three ways: through the C++ library, a command-line interface  
57 clvotess, and a Python wrapper interface pyvotess. The C++ library offers a simple  
58 interface with a tessellate function that computes the mesh. The Python wrapper, mirrors

the functionality of the C++ version, with native numpy array support, providing ease of use for Python-based workflows.

The behavior of votess can be fine-tuned with run time parameters in order to (optionally) optimize runtime performance.

## Acknowledgements

CB and DN acknowledge funding from the Deutsche Forschungsgemeinschaft (DFG) through an Emmy Noether Research Group (grant number NE 2441/1-1).

## References

- Bowyer, A. (1981). Computing dirichlet tessellations\*. *The Computer Journal*, 24(2), 162–166. <https://doi.org/10.1093/comjnl/24.2.162>
- Inria, P. A.-G. (2018). *Geogram: A programming library of geometric algorithms*. <http://alice.loria.fr/software/geogram/doc/html/index.html>
- Levy, B., Mohayaee, R., & Hausegger, S. von. (2021). A fast semidiscrete optimal transport algorithm for a unique reconstruction of the early universe. *Monthly Notices of the Royal Astronomical Society*, 506(1), 1165–1185. <https://doi.org/10.1093/mnras/stab1676>
- Lévy, B., & Bonneel, N. (2013). Variational anisotropic surface meshing with voronoi parallel linear enumeration. In X. Jiao & J.-C. Weill (Eds.), *Proceedings of the 21st international meshing roundtable* (pp. 349–366). Springer Berlin Heidelberg. ISBN: 978-3-642-33573-0
- Marot, C., Pellerin, J., & Remacle, J.-F. (2019). One machine, one minute, three billion tetrahedra. *International Journal for Numerical Methods in Engineering*, 117(9), 967–990. <https://doi.org/https://doi.org/10.1002/nme.5987>
- Ray, N., Sokolov, D., Lefebvre, S., & Lévy, B. (2018). Meshless voronoi on the GPU. *ACM Transactions on Graphics*, 37(6), 1–12. <https://doi.org/10.1145/3272127.3275092>
- Rycroft, C. (2009). *VORO++: A three-dimensional voronoi cell library in c++*. <https://doi.org/10.2172/946741>
- Springel, V. (2010). Moving-mesh hydrodynamics with the AREPO code. *Proceedings of the International Astronomical Union*, 6(S270), 203–206. <https://doi.org/10.1017/S1743921311000378>
- The CGAL Project. (2018). *CGAL user and reference manual* (4.12.1 ed.). CGAL Editorial Board. <https://doc.cgal.org/4.12.1/Manual/packages.html#PkgSpatialSortingSummary>
- Watson, D. F. (1981). Computing the n-dimensional delaunay tessellation with application to voronoi polytopes\*. *The Computer Journal*, 24(2), 167–172. <https://doi.org/10.1093/comjnl/24.2.167>
- Weygaert, R. van de. (1994). Fragmenting the universe. 3: The constructions and statistics of 3-d voronoi tessellations. *Astronomy and Astrophysics*, 283(2), 361–406.
- Wu, G., Tian, H., Lu, G., & Wang, W. (2023). ParVoro++: A scalable parallel algorithm for constructing 3D voronoi tessellations based on kd-tree decomposition. *Parallel Computing*, 115, 102995. <https://doi.org/10.1016/j.parco.2023.102995>